

TDDC78 - Programming of Parallel Computers

Lab4 - Particle Simulation using MPI

Authors:

David Tran - Davtr766

Jakob Bertlin - Jakbe457

I. INTRODUCTION

The fourth lab in the course TDDC78 introduces a particle simulation problem. The particle simulation works by instantiating particles in a box. The particles can either interact with other particles by colliding or walls.

In this assignment, the aim is to parallelize the particle simulation using MPI and verify the gas law:

$$pV = nRT$$

where p is the pressure, V is the volume (in this case area), n is the number of particles, R is a magic constant and T is the temperature, in this case the volume will be instead of area.

II. APPROACH

The approach to verify the gas law. It is hypothesized that the pressure should increase if the number of particles increases as n is directly proportional to p .

This report will also discuss the choice of distribution of particles between processors given the geometry domain of the implementation.

III. IMPLEMENTATION

This section covers the implementation of particles and the use of MPI.

A. Particles

The implementation was written in C++ and the particles was defined by a struct. The particles was managed by a vector of pairs, where each pair contains one particle and a boolean flag of whether the particle has collided with another particle/wall or not.

B. MPI

There are many ways to implement the MPI communication, however, in this report the approach was row-wise partitioning based on the amount of processors available. An illustration can be seen in figure 1.



Fig. 1: Simple illustration of the box area in the particle simulation. Each chunk is assigned to one processor where the horizontal length of the chunk depends on the number of processors and the vertical length is the same as the vertical size of the box.

Every processor instantiates the same amount of particles and the particles will only move in their respective chunk area. However, if particles move over to another chunk then they have to be sent to the processor that handles that specific chunk. Therefore, the implementation requires two vectors that acts as a buffer. One vector handles particles that have to be sent up and the other vector handles particles that have to be sent down.

The communication was done using *MPI_Isend* and *MPI_recv*. The implementation is written in a way such that particles are always sent up or down even though it may not contain any particles in the buffer. The buffers were not always the same size because for each time stamp, they may contain more particles or less particles depending on the simulation. Therefore, *MPI_Probe* was used

to determine how many particles were going to be received.

IV. RESULT

Figure 2 shows the dependence of number of cores with elapsed time and the pressure.

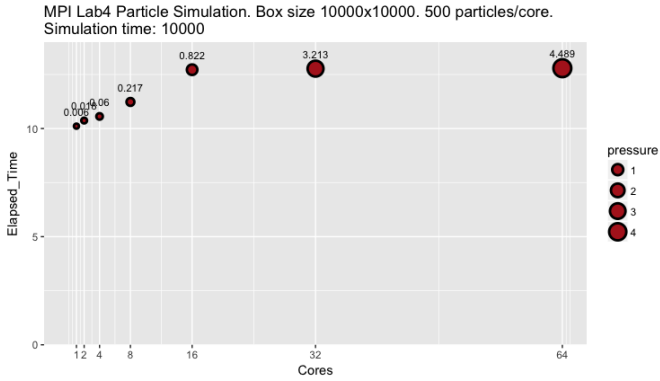


Fig. 2: Scatter plot where the x-axis are the number of cores, the y-axis is the elapsed time and the size of the points are determined by the amount of pressure.

V. DISCUSSION

The scatter plot in section IV bolsters previous hypothesis about the gas law.

Figure 2 verifies the gas law described in section II. The pressure increases as the number of particles increases. This proves that the pressure is proportional to the number of particles. The plot also shows that the elapsed time is more or less the same. The reason for that is each processor works on the same amount of particles.

The communication of particles between processors was implemented in a way that two send operations are required, either up or down. The advantage over having a grid is that row-wise partitioning allows less communication however each communication are larger. However, in our case, less communication is better than more communication because the cost of initializing another communication operation is more expensive than sending a larger one.

APPENDIX

```

1 #include <cstdlib>
2 #include <ctime>
3 #include <cstdio>
4 #include <cmath>
5 #include <limits>
6 #include <iostream>
7 #include <vector>
8 #include <mpi.h>
9 #include <algorithm>
10 #include <utility>
11 #include "coordinate.h"
12 #include "definitions.h"
13 #include "physics.h"
14
15
16 //Feel free to change this program to facilitate parallelization.
17
18 float randl(){
19     return (rand()/(float) RAND_MAX);
20 }
21
22 int calcRowRank(float y){
23     int n_proc;
24     MPI_Comm_size(MPI_COMM_WORLD, &n_proc);
25     float rowSplit = (float)BOX_VERT_SIZE / n_proc;
26     return (int)(y/rowSplit);
27 }
28
29 bool boundaryCheck(Particle &p, cord_t wall){
30
31     if((p.y < wall.y0 || p.y > wall.y1) && (p.y > 0 && p.y < BOX_VERT_SIZE) )
32         return true;
33     else
34         return false;
35 }
36
37
38 int main(int argc, char** argv){
39     /* Define variables */
40     unsigned int time_stamp = 0, time_max;
41     float pressure = 0;
42     int rank(), world(), root{0};
43     std::vector<std::pair<Particle, bool>> Particles;
44     std::vector<Particle> sendParticlesUp;
45     std::vector<Particle> sendParticlesDown;
46     cord_t wall;
47     Particle *recbuf;
48     double start_time{0}, end_time{0};
49
50
51     /* Initialize MPI environment */
52     MPI_Init(nullptr, nullptr);
53     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
54     MPI_Comm_size(MPI_COMM_WORLD, &world);
55
56     /* Create MPI data type for particles */
57     MPI_Datatype MPI_Particle, oldtypes[1];
58     int blockcounts[1];
59     MPI_Aint offsets[1];
60     offsets[0] = 0;
61     oldtypes[0] = MPI_FLOAT;
62     blockcounts[0] = 4;
63     MPI_Type_create_struct(1, blockcounts, offsets, oldtypes, &MPI_Particle);
64     MPI_Type_commit(&MPI_Particle);
65     long localsum{0}, globalsum{0}, max_sent{0}, min_sent{std::numeric_limits<long>::max()}, global_max{0}, global_min{0};
66
67
68     // parse arguments
69     if(argc != 2) {
70         fprintf(stderr, "Usage: %s simulation_time\n", argv[0]);
71         fprintf(stderr, "For example: %s 10\n", argv[0]);
72         exit(1);
73     }
74
75     time_max = atoi(argv[1]);
76
77     //std::cout << "Rank " << rank << " out of " << world << "\n";
78
79     //Initializes the wall and Particle on the root/master processor.
80
81     // 1. set the walls
82     float rowSplit = (float)(BOX_VERT_SIZE/world);
83     wall.y0 = rank * rowSplit;
84     wall.y1 = wall.y0 + rowSplit;
85     wall.x0 = 0;
86     wall.x1 = BOX_HORIZ_SIZE;
87
88
89     // 2. allocate particle buffer and initialize the Particle
90     Particles = std::vector<std::pair<Particle, bool>>(INIT_NO_PARTICLES);
91
92
93     srand( time(nullptr) + 1234 );
94
95     float r, a;
96     for(int i=0; i<INIT_NO_PARTICLES; i++){

```

```

97 // initialize random position
98 Particles[i].first.x = static_cast<float>(wall.x0 + rand1()*BOX_HORIZ_SIZE);
99 Particles[i].first.y = (wall.y0 + rand1()*rowSplit);
100
101 // initialize random velocity
102 r = rand1()*MAX_INITIAL_VELOCITY;
103 if(r > 50) r = 49;
104 a = static_cast<float>(rand1()*2*PI);
105 Particles[i].first.vx = r*cos(a);
106 Particles[i].first.vy = r*sin(a);
107 }
108 MPI_Request req[2];
109 MPI_Status statUp, statDown, statRec;
110 Particle *sendBufDown = new Particle[1], *sendBufUp = new Particle[1];
111 int sendcountDown, sendcountUp;
112
113
114 start_time = MPI_Wtime(); //start MPI::Wtime();
115 /* Main loop */
116 for (time_stamp=0; time_stamp<time_max; time_stamp++) { // for each time stamp
117
118
119 /* Initialize values */
120 for(auto& p : Particles)
121     p.second = false;
122
123 sendParticlesUp.clear();
124 sendParticlesDown.clear();
125
126 /* Main collision loop */
127 for(auto p = Particles.begin(); p != Particles.end(); ++p) { // for all Particle
128     if(p->second) continue;
129
130 /* check for collisions */
131 for(auto pp=p+1; pp != Particles.end(); ++pp){
132     if(pp->second) continue;
133
134     float t=collide(&(p->first), &(pp->first));
135     if(t!=-1){ // collision
136         p->second = pp->second = true;
137         interact(&(p->first), &(pp->first), t);
138
139         /******
140          * For LAB 5 TOTALVIEW Debug *
141          *****/
142         /*if(boundaryCheck((p->first), wall)) {
143             if(calcRowRank(p->first.y) < rank)
144                 sendParticlesUp.emplace_back(p->first);
145             else
146                 sendParticlesDown.emplace_back(p->first);
147
148             //std::cout << "Erasing... \n" ;
149             Particles.erase( p );
150         }
151         if(boundaryCheck(pp->first, wall)) {
152             if(calcRowRank(pp->first.y) < rank)
153                 sendParticlesUp.emplace_back(pp->first);
154             else
155                 sendParticlesDown.emplace_back(pp->first);
156
157             //std::cout << "Erasing... \n" ;
158             Particles.erase( pp );
159         }
160         */
161
162         break; // only check collision of two Particle
163     }
164 }
165
166 }
167
168
169 // move Particle that has not collided with another
170 for(auto p = Particles.begin(); p != Particles.end(); ) {
171     if(!p->second){
172         feuler(&(p->first), 1);
173         pressure += wall_collide(&(p->first), wall);
174     }
175
176     if(boundaryCheck(p->first, wall)) {
177         if (calcRowRank(p->first.y) < rank)
178             sendParticlesUp.emplace_back(p->first);
179         else
180             sendParticlesDown.emplace_back(p->first);
181
182         p = Particles.erase(p);
183     }
184     else{
185         ++p;
186     }
187 }
188
189
190
191 // Send to another process
192 if(rank != root) {
193     //std::cout << "Sending " << sendParticlesUp.size() << " from rank " << rank << "... \n" ;
194     delete[] sendBufUp;

```

```

195     sendcountUp = static_cast<int>(sendParticlesUp.size());
196     sendBufUp = new Particle[sendcountUp];
197     max_sent = sendcountUp > max_sent ? sendcountUp : max_sent;
198     min_sent = sendcountUp < min_sent ? sendcountUp : min_sent;
199     localsum += sendcountUp;
200     for(int i=0; i<sendParticlesUp.size(); i++)
201         sendBufUp[i] = sendParticlesUp[i];
202
203     MPI_Isend(sendBufUp, sendcountUp, MPI_Particle, rank-1, 0, MPI_COMM_WORLD, &(req[0]));
204 }
205
206
207 if(rank != world-1) {
208     //std::cout << "Sending " << sendParticlesDown.size() << " from rank " << rank << "... \n" ;
209     delete[] sendBufDown;
210     sendcountDown = static_cast<int>(sendParticlesDown.size());
211     sendBufDown = new Particle[sendcountDown];
212     max_sent = sendcountDown > max_sent ? sendcountDown : max_sent;
213     min_sent = sendcountDown < min_sent ? sendcountDown : min_sent;
214     localsum += sendcountDown;
215     for(int i=0; i<sendParticlesDown.size(); i++)
216         sendBufDown[i] = sendParticlesDown[i];
217
218     MPI_Isend(sendBufDown, sendcountDown, MPI_Particle, rank+1, 0, MPI_COMM_WORLD, &(req[1]));
219 }
220
221
222 if(rank != root){
223     MPI_Probe(rank-1, MPI_ANY_TAG, MPI_COMM_WORLD, &statRec);
224     int recCount{};
225     MPI_Get_count(&statRec, MPI_Particle, &recCount);
226     //std::cout << "Received " << recCount << " from rank: " << statRec.MPI_SOURCE << "\n";
227     if(recCount != 0)
228         recbuf = new Particle[recCount];
229     MPI_Recv(recbuf, recCount, MPI_Particle, statRec.MPI_SOURCE, statRec.MPI_TAG, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
230     if(recCount != 0) {
231         for (int i = 0; i < recCount; i++)
232             Particles.emplace_back(std::make_pair(recbuf[i], false));
233
234         delete[] recbuf;
235     }
236 }
237
238
239 if(rank != world-1){
240     MPI_Probe(rank+1, MPI_ANY_TAG, MPI_COMM_WORLD, &statRec);
241     int recCount{};
242     MPI_Get_count(&statRec, MPI_Particle, &recCount);
243     //std::cout << "Received " << recCount << " from rank: " << statRec.MPI_SOURCE << "\n";
244     if(recCount != 0)
245         recbuf = new Particle[recCount];
246     MPI_Recv(recbuf, recCount, MPI_Particle, statRec.MPI_SOURCE, statRec.MPI_TAG, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
247     if(recCount != 0) {
248         for (int i = 0; i < recCount; i++)
249             Particles.emplace_back(std::make_pair(recbuf[i], false));
250
251         delete[] recbuf;
252     }
253 }
254
255 MPI_Barrier(MPI_COMM_WORLD);
256
257 }
258
259
260 // Get total pressure from all processes
261 float totalpress = 0;
262 int n = (int)Particles.size(), totalp = 0;
263
264 //std::cout << localsum << std::endl;
265 MPI_Reduce(&pressure, &totalpress, 1, MPI_FLOAT, MPI_SUM, 0, MPI_COMM_WORLD);
266 MPI_Reduce(&localsum, &globalsum, 1, MPI_LONG, MPI_SUM, 0, MPI_COMM_WORLD);
267 MPI_Reduce(&min_sent, &global_min, 1, MPI_LONG, MPI_MIN, 0, MPI_COMM_WORLD);
268 MPI_Reduce(&max_sent, &global_max, 1, MPI_LONG, MPI_MAX, 0, MPI_COMM_WORLD);
269 MPI_Reduce(&n, &totalp, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);
270
271
272 end_time = MPI_Wtime(); //start MPI::Wtime();
273 if(rank == root){
274     std::cout << "Size of box: " << BOX_HORIZ_SIZE << "x" << BOX_VERT_SIZE << "\n";
275     std::cout << "Total particles " << totalp << "\n";
276     printf("Average pressure = %f\n", totalpress / (WALL_LENGTH*time_max));
277     printf("Elapsed time = %f seconds\n", end_time-start_time);
278     std::cout << "Average sent particles = " << globalsum/time_max << "\n";
279     std::cout << "Minimum particles sent = " << global_min << "\n";
280     std::cout << "Maximum particles sent = " << global_max << "\n";
281 }
282
283
284 MPI_Finalize();
285 return 0;
286 }

```